

ML Posledni Ukol

Matyas Vondra 22.6. 2023

Kapitoly:

1. Priprava dat
2. Referencni jednoduchy linearni regrese
3. Elastic net regrese
4. Rozhodovaci stromy
5. Support Vector Regression
6. Neuronova sit
7. Zaver

Priprava dat

```
In [1]: import numpy as np  
import pandas as pd
```

```
In [2]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [3]: data =  
pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Workshop/Clean_Dataset.c
```

```
In [4]: data
```

Out [4]:	Unnamed: 0	airline	flight	source_city	departure_time	stops	arrival_time	dest
0	0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	
1	1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	
2	2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	
3	3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	
4	4	Vistara	UK-963	Delhi	Morning	zero	Morning	
...
300148	300148	Vistara	UK-822	Chennai	Morning	one	Evening	
300149	300149	Vistara	UK-826	Chennai	Afternoon	one	Night	
300150	300150	Vistara	UK-832	Chennai	Early_Morning	one	Night	
300151	300151	Vistara	UK-828	Chennai	Early_Morning	one	Evening	
300152	300152	Vistara	UK-822	Chennai	Morning	one	Evening	

300153 rows x 12 columns

```
In [5]: # vyber pouze radku business tridy
business = data.loc[data['class'] == "Business"]
```

```
In [6]: business.shape
```

Out[6]: (93487, 12)

```
In [7]: business
```

Out [7]:

	Unnamed: 0	airline	flight	source_city	departure_time	stops	arrival_time	desti
206666	206666	Air_India	AI-868	Delhi	Evening	zero	Evening	
206667	206667	Air_India	AI-624	Delhi	Evening	zero	Night	
206668	206668	Air_India	AI-531	Delhi	Evening	one	Night	
206669	206669	Air_India	AI-839	Delhi	Night	one	Night	
206670	206670	Air_India	AI-544	Delhi	Evening	one	Night	
...
300148	300148	Vistara	UK-822	Chennai	Morning	one	Evening	
300149	300149	Vistara	UK-826	Chennai	Afternoon	one	Night	
300150	300150	Vistara	UK-832	Chennai	Early_Morning	one	Night	
300151	300151	Vistara	UK-828	Chennai	Early_Morning	one	Evening	
300152	300152	Vistara	UK-822	Chennai	Morning	one	Evening	

93487 rows x 12 columns

```
In [8]: # odebrani zbytecznych sloupcu
business = business.drop('Unnamed: 0', axis=1)
business = business.drop('class', axis=1)
```

```
In [9]: business
```

Out [9]:

	airline	flight	source_city	departure_time	stops	arrival_time	destination_city
206666	Air_India	AI-868	Delhi	Evening	zero	Evening	Mumbai
206667	Air_India	AI-624	Delhi	Evening	zero	Night	Mumbai
206668	Air_India	AI-531	Delhi	Evening	one	Night	Mumbai
206669	Air_India	AI-839	Delhi	Night	one	Night	Mumbai
206670	Air_India	AI-544	Delhi	Evening	one	Night	Mumbai
...
300148	Vistara	UK-822	Chennai	Morning	one	Evening	Hyderabad
300149	Vistara	UK-826	Chennai	Afternoon	one	Night	Hyderabad
300150	Vistara	UK-832	Chennai	Early_Morning	one	Night	Hyderabad
300151	Vistara	UK-828	Chennai	Early_Morning	one	Evening	Hyderabad
300152	Vistara	UK-822	Chennai	Morning	one	Evening	Hyderabad

93487 rows × 10 columns

Kontrola typu a hodnot promenných

```
In [10]: business.dtypes
```

```
Out[10]: airline          object
flight          object
source_city     object
departure_time  object
stops           object
arrival_time    object
destination_city object
duration        float64
days_left      int64
price           int64
dtype: object
```

```
In [11]: business.departure_time.unique()
```

```
Out[11]: array(['Evening', 'Night', 'Early_Morning', 'Morning', 'Afternoon',
                'Late_Night'], dtype=object)
```

```
In [12]: business.stops.unique()
```

```
Out[12]: array(['zero', 'one', 'two_or_more'], dtype=object)
```

```
In [13]: business.arrival_time.unique()
```

```
Out[13]: array(['Evening', 'Night', 'Afternoon', 'Morning', 'Late_Night',  
               'Early_Morning'], dtype=object)
```

```
In [14]: business.source_city.unique()
```

```
Out[14]: array(['Delhi', 'Mumbai', 'Bangalore', 'Kolkata', 'Hyderabad', 'Chennai'],  
               dtype=object)
```

```
In [15]: business.destination_city.unique()
```

```
Out[15]: array(['Mumbai', 'Bangalore', 'Kolkata', 'Hyderabad', 'Chennai', 'Delhi'],  
               dtype=object)
```

```
In [16]: business.airline.unique()
```

```
Out[16]: array(['Air_India', 'Vistara'], dtype=object)
```

```
In [17]: business.flight.unique()
```

```
Out[17]: array(['AI-868', 'AI-624', 'AI-531', 'AI-839', 'AI-544', 'UK-985',
               'AI-479', 'AI-473', 'UK-871', 'UK-977', 'AI-504', 'AI-807',
               'AI-540', 'AI-537', 'UK-817', 'AI-762', 'AI-764', 'UK-707',
               'UK-809', 'UK-813', 'UK-837', 'UK-953', 'AI-887', 'AI-665',
               'AI-805', 'AI-678', 'AI-636', 'AI-441', 'AI-435', 'UK-683',
               'AI-403', 'UK-963', 'UK-955', 'AI-411', 'AI-811', 'AI-453',
               'UK-637', 'UK-927', 'AI-483', 'AI-542', 'AI-560', 'AI-406',
               'UK-829', 'UK-879', 'UK-899', 'UK-627', 'UK-945', 'AI-451',
               'AI-885', 'AI-877', 'UK-995', 'UK-673', 'AI-429', 'UK-859',
               'AI-439', 'UK-839', 'UK-833', 'UK-819', 'UK-801', 'UK-815',
               'UK-706', 'AI-465', 'AI-512', 'AI-767', 'AI-401', 'UK-835',
               'UK-737', 'UK-847', 'UK-951', 'AI-475', 'AI-431', 'UK-933',
               'AI-502', 'AI-506', 'AI-803', 'UK-855', 'UK-705', 'AI-499',
               'AI-471', 'AI-485', 'AI-423', 'AI-415', 'UK-747', 'AI-481',
               'UK-981', 'UK-975', 'UK-993', 'UK-943', 'UK-941', 'UK-811',
               'UK-807', 'AI-837', 'AI-865', 'UK-727', 'AI-491', 'AI-883',
               'AI-437', 'AI-861', 'AI-641', 'UK-812', 'UK-671', 'AI-433',
               'AI-487', 'AI-407', 'AI-889', 'AI-801', 'UK-641', 'AI-895',
               'AI-879', 'AI-489', 'AI-891', 'AI-493', 'AI-888', 'AI-867',
               'AI-631', 'AI-637', 'UK-996', 'UK-940', 'AI-619', 'UK-875',
               'UK-988', 'AI-681', 'AI-774', 'UK-865', 'UK-823', 'AI-806',
               'AI-809', 'AI-687', 'AI-660', 'AI-635', 'AI-442', 'UK-958',
               'UK-928', 'UK-960', 'AI-625', 'UK-613', 'AI-649', 'UK-970',
               'UK-877', 'UK-621', 'UK-944', 'AI-685', 'AI-683', 'AI-669',
               'UK-910', 'UK-651', 'AI-679', 'AI-639', 'AI-570', 'UK-863',
               'UK-851', 'UK-825', 'UK-853', 'AI-671', 'UK-775', 'UK-771',
               'AI-655', 'UK-655', 'AI-645', 'AI-651', 'AI-607', 'AI-673',
               'UK-994', 'AI-643', 'AI-611', 'UK-773', 'UK-954', 'UK-950',
               'UK-902', 'UK-653', 'UK-873', 'UK-986', 'UK-821', 'UK-827',
               'UK-845', 'UK-930', 'UK-849', 'UK-841', 'UK-857', 'AI-864',
               'UK-861', 'AI-615', 'AI-695', 'AI-657', 'AI-633', 'AI-652',
               'AI-629', 'AI-627', 'AI-601', 'AI-623', 'AI-603', 'UK-818',
               'AI-808', 'UK-850', 'UK-808', 'UK-820', 'UK-802', 'AI-505',
               'AI-501', 'AI-804', 'AI-503', 'UK-810', 'AI-507', 'AI-640',
               'UK-657', 'UK-816', 'UK-854', 'UK-852', 'UK-858', 'AI-573',
               'AI-610', 'AI-565', 'UK-867', 'AI-738', 'AI-516', 'AI-523',
               'AI-748', 'UK-866', 'UK-864', 'UK-814', 'UK-846', 'AI-776',
               'AI-604', 'AI-772', 'AI-564', 'AI-583', 'UK-774', 'AI-763',
               'AI-402', 'AI-768', 'AI-770', 'UK-738', 'UK-778', 'AI-729',
               'AI-781', 'AI-526', 'AI-713', 'AI-765', 'AI-773', 'AI-787',
               'UK-776', 'UK-772', 'AI-721', 'UK-720', 'AI-745', 'AI-732',
               'AI-715', 'AI-747', 'AI-424', 'UK-708', 'AI-743', 'AI-780',
               'AI-775', 'AI-771', 'AI-840', 'AI-541', 'UK-860', 'UK-880',
               'UK-830', 'AI-559', 'AI-543', 'UK-890', 'AI-698', 'AI-508',
               'AI-525', 'UK-870', 'UK-876', 'UK-878', 'AI-838', 'AI-515',
               'UK-874', 'AI-616', 'AI-420', 'AI-620', 'AI-546', 'UK-832',
               'UK-838', 'AI-440', 'AI-539', 'AI-538', 'AI-430', 'UK-836',
               'AI-569', 'AI-672', 'AI-766', 'UK-834', 'UK-822', 'AI-549',
               'UK-824', 'UK-828', 'UK-826', 'AI-545', 'AI-551', 'AI-563',
               'AI-509'], dtype=object)
```

One hot encoding promennych

```
In [18]: transformed = pd.get_dummies(business)
```

In [19]: transformed

Out[19]:

	duration	days_left	price	airline_Air_India	airline_Vistara	flight_AI-401	flight_AI-402
206666	2.00	1	25612	1	0	0	0
206667	2.25	1	25612	1	0	0	0
206668	24.75	1	42220	1	0	0	0
206669	26.50	1	44450	1	0	0	0
206670	6.67	1	46690	1	0	0	0
...
300148	10.08	49	69265	0	1	0	0
300149	10.42	49	77105	0	1	0	0
300150	13.83	49	79099	0	1	0	0
300151	10.00	49	81585	0	1	0	0
300152	10.08	49	81585	0	1	0	0

93487 rows x 327 columns

In [20]: *# Normalizace dat (nakonec nepouzita)*
normalized_df =
(transformed-transformed.min())/(transformed.max()-transformed.min())

In [21]: normalized_df

Out[21]:

	duration	days_left	price	airline_Air_India	airline_Vistara	flight_AI-401	flight_AI-402
206666	0.021390	0.0	0.122552	1.0	0.0	0.0	0
206667	0.026738	0.0	0.122552	1.0	0.0	0.0	0
206668	0.508021	0.0	0.272078	1.0	0.0	0.0	0
206669	0.545455	0.0	0.292155	1.0	0.0	0.0	0
206670	0.121283	0.0	0.312323	1.0	0.0	0.0	0
...
300148	0.194225	1.0	0.515571	0.0	1.0	0.0	0
300149	0.201497	1.0	0.586157	0.0	1.0	0.0	0
300150	0.274439	1.0	0.604109	0.0	1.0	0.0	0
300151	0.192513	1.0	0.626491	0.0	1.0	0.0	0
300152	0.194225	1.0	0.626491	0.0	1.0	0.0	0

93487 rows x 327 columns

```
In [22]: # Rozdeleni na train a test
from sklearn.model_selection import train_test_split
X = np.array(transformed.loc[:, transformed.columns != "price"])
y = np.array(transformed["price"])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
                                                    random_state = 1)
```

```
In [23]: X_train
```

```
Out[23]: array([[19.75, 19. , 1. , ..., 0. , 1. , 0. ],
               [12.17, 5. , 1. , ..., 0. , 0. , 0. ],
               [22.08, 10. , 0. , ..., 0. , 0. , 0. ],
               ...,
               [14.92, 14. , 0. , ..., 1. , 0. , 0. ],
               [24.33, 48. , 1. , ..., 0. , 0. , 1. ],
               [26.17, 30. , 0. , ..., 0. , 1. , 0. ]])
```

```
In [27]: X_train.shape
```

```
Out[27]: (74789, 326)
```

```
In [24]: y_train
```

```
Out[24]: array([57405, 43729, 57992, ..., 56588, 59033, 58394])
```

Definice ztratove funkce. Vybral jsem si MAPE, primarne pro její snadnou interpretovatelnost.

```
In [25]: def mape(y, yhat):
          return np.mean([np.abs(i-j)/j for i,j in zip(yhat, y)])
```

Referencni jednoduchá lineární regrese

```
In [26]: from sklearn.linear_model import LinearRegression
```

```
In [27]: linear_regression = LinearRegression()
linear_regression.fit(X_train, y_train)
```

```
Out[27]: ▼ LinearRegression
LinearRegression()
```

```
In [28]: linear_regression.score(X_test, y_test) # out of sample index determinace
```

```
Out[28]: -330096676127365.44
```

```
In [29]: yhat = linear_regression.predict(X_test)
```

```
In [30]: # mape
mape(y_test, yhat)
```


Out [30]: 58122.98456233965

Dle očekávání měla úplně nejednodušší lineární regrese přinést výsledky :D

Elastic net regrese

Elastic net jsem se rozhodl použít jelikož data neznám tak dobře, abych mohl určit, zda je lepší L1 nebo L2 regularizace. Jelikož je vysoký počet proměnných, očekávám multikolinearitě.

```
In [31]: from sklearn.linear_model import ElasticNet
```

```
In [32]: # základní model
el = ElasticNet(alpha=1.0, l1_ratio=0.5)

el.fit(X_train, y_train)

mape(y_test, el.predict(X_test))
```

Out [32]: 0.19153055967999336

Už tenhle výsledek je za mě velmi dobrý. Průměrná out-of-sample absolutní chyba pod 20% je slušná :D


```
In [33]: # hledání optimálního lambdy/alpha
lambdas = list()
for i in range(1000, 10_000, 100):
    res = ElasticNet(alpha = i/10_000, l1_ratio=0.95) # 0.95 doplneno zpětně
    res.fit(X_train, y_train)
    lambdas.append(mape(y_train, res.predict(X_train))) # insample mape
```

```
In [34]: # hledání optimálního parametru pro volbu L1/L2 regularizace
l1_wts = list()
for i in range(1000, 10_000, 100):
    res = ElasticNet(alpha = 0.1, l1_ratio = i/10000)
    res.fit(X_train, y_train)
    l1_wts.append(mape(y_train, res.predict(X_train)))
```

```
In [35]: import matplotlib.pyplot as plt
```


```
In [36]: plt.plot(range(90), lambdas)
# zda se ze lambda bude velmi malá, asi 0,1
```

Out [36]: [<matplotlib.lines.Line2D at 0x7fb3a0b43820>]



```
In [37]: plt.plot(range(90), l1_wts)
# zda se ze vhodnejši je lasso regrese (vyšši hodnota lambda)
```

```
Out[37]: [<matplotlib.lines.Line2D at 0x7fb3a099c4c0>]
```



Nakonec se jako vhodný model jeví lineární regrese s L1 regularizací (Lasso), která je

schopna vynechavat nektère parametry. Hodnota α/λ vychází optimalne jako pomerne nizka, tedy nedochází k tak velké penalizaci. MAPE vychází ~11%, což je podle mě velmi dobře

```
In [38]: el = ElasticNet(alpha=0.1, l1_ratio=0.95)

el.fit(X_train, y_train)

mape(y_test, el.predict(X_test))
```

Out[38]: 0.1133152803325097

Rozhodovací stromy

Základní regresní strom

Použit Randomized Search (podobný jako grid, ale netestuje všechny kombinace, jen náhodný výběr)

```
In [39]: from sklearn.tree import DecisionTreeRegressor
```

```
In [40]: dtr = DecisionTreeRegressor()
```

```
In [41]: random_grid = {
    "ccp_alpha": np.linspace(0.0001, 0.1, 10), # 10 různých ccp
    "criterion": ["squared_error", "friedman_mse", "absolute_error", "poisson_deviance"],
    "min_samples_leaf": [1, 2, 5, 10]
}
```

```
In [42]: from sklearn.model_selection import RandomizedSearchCV
```

```
In [43]: randomizerCV = RandomizedSearchCV(estimator = dtr,
    param_distributions = random_grid,
    n_iter = 3,
    cv=2, verbose=2,
    random_state = 1)
```

```
In [44]: randomizerCV.fit(X_train, y_train)
```

```
Fitting 2 folds for each of 3 candidates, totalling 6 fits
[CV] END ccp_alpha=0.0112, criterion=poisson, min_samples_leaf=2; total time= 1.4s
[CV] END ccp_alpha=0.0112, criterion=poisson, min_samples_leaf=2; total time= 1.4s
[CV] END ccp_alpha=0.0223, criterion=absolute_error, min_samples_leaf=5; total time= 6.7min
[CV] END ccp_alpha=0.0223, criterion=absolute_error, min_samples_leaf=5; total time= 6.5min
[CV] END ccp_alpha=0.0001, criterion=poisson, min_samples_leaf=5; total time= 0.8s
[CV] END ccp_alpha=0.0001, criterion=poisson, min_samples_leaf=5; total time= 0.8s
```

```
Out[44]: RandomizedSearchCV
  estimator: DecisionTreeRegressor
    DecisionTreeRegressor
```

```
In [45]: # nalezeni optimalnich parametru pro nas strom
randomizerCV.best_params_
```

```
Out[45]: {'min_samples_leaf': 5, 'criterion': 'poisson', 'ccp_alpha': 0.0001}
```

```
In [46]: regression_tree = DecisionTreeRegressor(min_samples_leaf = 5,
                                                criterion = 'poisson',
                                                ccp_alpha = 0.1)
```

```
In [47]: regression_tree.fit(X_train, y_train)
```

```
Out[47]: DecisionTreeRegressor
DecisionTreeRegressor(ccp_alpha=0.1, criterion='poisson', min_samples_leaf=5)
```

```
In [48]: # velmi nizka out-of-sample chyba
mape(y_test, regression_tree.predict(X_test))
```

```
Out[48]: 0.0429253696015026
```

Bagging

100 bootstrapovanych (nahodny vyber pozorovani s vracenim) regresnich stromu.
Vysledna chyba vznikne zprumerovanim.

```
In [49]: from sklearn.ensemble import BaggingRegressor
```

```
In [50]: bagging = BaggingRegressor(n_estimators=100, random_state=1)
```

```
In [51]: bagging.fit(X_train, y_train)
```

```
Out[51]: ▼ BaggingRegressor
          BaggingRegressor(n_estimators=100, random_state=1)
```

```
In [52]: mape(y_test, bagging.predict(X_test))
```

```
Out[52]: 0.02806644129803859
```

Velmi dobry vysledek, zatim nejlepsi. Mozna byl nas puvodni strom trochu pre-uceny, coz se nam podarilo baggingem vyresit.

Random forest

Opet 100 nahodnych stromu, vctne nahodneho vyberu promennych. Pouzity stejne parametry jako pri klasickem stromu.

```
In [53]: from sklearn.ensemble import RandomForestRegressor
```

[illegible]

```
In [55]: random_forest.fit(X_train, y_train)
```

```
Out[55]: ▼ RandomForestRegressor
RandomForestRegressor(ccp_alpha=0.1, criterion='poisson', min_samples_leaf=5,
                      random_state=1)
```

```
In [56]: mape(y_test, random_forest.predict(X_test))
```

```
Out[56]: 0.04078963466466962
```

Random forest ma sice mensi chybu nez puvodni strom, ale neni tak dobry jako obycejny Bagging. Pravdepodobne by pomohlo testovat vice ruznych settingu random forestu a najit nejaky lepsi.

Boosting

Predikce chyby predchoziho stromu. Metoda: gradient.

```
In [57]: from sklearn.ensemble import HistGradientBoostingRegressor
```

[illegible]

```
min_samples_leaf = 5)
```

```
In [59]: boosting.fit(X_train, y_train)
```

```
Out[59]: ▾ HistGradientBoostingRegressor  
HistGradientBoostingRegressor(loss='poisson', min_samples_leaf=5)
```

```
In [60]: mape(y_test, boosting.predict(X_test))
```

```
Out[60]: 0.06874158115299706
```

Zde si takto "vysokou" chybu nedokazu uplne vysvetlit. Mozna zkosit upravit learning rate (default je 0.1) nebo treba L2 regularizaci, kterou model nabizi.

Support Vector Regression

Spise pro zajimavost, nemam tolik zkusenosti s vyberem parametru pro kernely a trenink trval velmi, velmi dlouho.

Jinak pokud tomu dobre rozumim, tak SVR je vlastne celkem podobna klasicke regresi. Diky pouziti kernelu je ale mozne dobre fitovat model i na nelinearni data.

```
In [61]: from sklearn.svm import SVR
```

```
In [62]: kernels = ["linear", "poly", "rbf"]
```

```
In [63]: for kernel in kernels:  
    model = SVR(kernel = kernel)  
    model.fit(X_train, y_train)  
    yhat = model.predict(X_test)  
    print(f"{kernel}: {mape(y_test, yhat): .4f}")
```

```
linear: 0.1671
```

```
poly: 0.2308
```

```
rbf: 0.2304
```

Jak jiz mozna bylo mozne videt drive, nase data jsou dobre rozdelitelna i linearne (dobre vysledky regularizovane linearni regrese). Tim si vysvetluji horsi skore SVR s pouzitim nelinearnich kernelu nez kernelu `linear`.

Neuronova sit

```
In [29]: from keras.models import Sequential  
from keras.layers import Dense, Dropout, BatchNormalization
```

Neuronka s jednou skrytou vrstvou

- Input layer = 326 neuronu
- Hidden layer = 16 neuronu
- Batch normalizace vstupu
- Dropout 20% proti over fittingu
- Output layer = 1 neuron

Vsechny aktivace RELU

Batch size 16

```
In [30]: # Vytvoření a kompilace modelu
model = Sequential()
model.add(Dense(16, input_shape=(326,), activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.2))
model.add(Dense(1, activation='relu'))

model.compile(loss='mean_squared_error', optimizer='adam',
              metrics='mean_absolute_percentage_error')

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 16)	5232
batch_normalization (Batch Normalization)	(None, 16)	64
dropout (Dropout)	(None, 16)	0
dense_2 (Dense)	(None, 1)	17

```
=====
Total params: 5,313
Trainable params: 5,281
Non-trainable params: 32
=====
```

```
In [35]: model.fit(X_train, y_train,
                  epochs=20, batch_size=16,
                  verbose=1)
```

Epoch 1/20
4675/4675 [=====] - 14s 3ms/step - loss: 10318147
2.0000 - mean_absolute_percentage_error: 15.3554
Epoch 2/20
4675/4675 [=====] - 9s 2ms/step - loss: 100858056.
0000 - mean_absolute_percentage_error: 15.0977
Epoch 3/20
4675/4675 [=====] - 11s 2ms/step - loss: 99419680.
0000 - mean_absolute_percentage_error: 14.9246
Epoch 4/20
4675/4675 [=====] - 11s 2ms/step - loss: 99728944.
0000 - mean_absolute_percentage_error: 14.9564
Epoch 5/20
4675/4675 [=====] - 10s 2ms/step - loss: 98279192.
0000 - mean_absolute_percentage_error: 14.8108
Epoch 6/20
4675/4675 [=====] - 10s 2ms/step - loss: 96285328.
0000 - mean_absolute_percentage_error: 14.7581
Epoch 7/20
4675/4675 [=====] - 10s 2ms/step - loss: 96448688.
0000 - mean_absolute_percentage_error: 14.6732
Epoch 8/20
4675/4675 [=====] - 10s 2ms/step - loss: 96060464.
0000 - mean_absolute_percentage_error: 14.7116
Epoch 9/20
4675/4675 [=====] - 10s 2ms/step - loss: 94731232.
0000 - mean_absolute_percentage_error: 14.5567
Epoch 10/20
4675/4675 [=====] - 11s 2ms/step - loss: 95507864.
0000 - mean_absolute_percentage_error: 14.6669
Epoch 11/20
4675/4675 [=====] - 11s 2ms/step - loss: 94675704.
0000 - mean_absolute_percentage_error: 14.6384
Epoch 12/20
4675/4675 [=====] - 9s 2ms/step - loss: 94830560.0
000 - mean_absolute_percentage_error: 14.6619
Epoch 13/20
4675/4675 [=====] - 11s 2ms/step - loss: 94147432.
0000 - mean_absolute_percentage_error: 14.6477
Epoch 14/20
4675/4675 [=====] - 11s 2ms/step - loss: 94580928.
0000 - mean_absolute_percentage_error: 14.7146
Epoch 15/20
4675/4675 [=====] - 11s 2ms/step - loss: 94750824.
0000 - mean_absolute_percentage_error: 14.6743
Epoch 16/20
4675/4675 [=====] - 10s 2ms/step - loss: 92805576.
0000 - mean_absolute_percentage_error: 14.5660
Epoch 17/20
4675/4675 [=====] - 11s 2ms/step - loss: 93521728.
0000 - mean_absolute_percentage_error: 14.5580
Epoch 18/20
4675/4675 [=====] - 11s 2ms/step - loss: 93022568.
0000 - mean_absolute_percentage_error: 14.6101
Epoch 19/20
4675/4675 [=====] - 11s 2ms/step - loss: 93330144.


```
0000 - mean_absolute_percentage_error: 14.5814
Epoch 20/20
4675/4675 [=====] - 11s 2ms/step - loss: 93659648.
0000 - mean_absolute_percentage_error: 14.6387
```

Out[35]: <keras.callbacks.History at 0x7f7539652da0>

```
In [47]: # in-sample. Je mozne, ze keras nasobi MAPE *2?
mape(y_train, model.predict(X_train))
```

```
2338/2338 [=====] - 5s 2ms/step
```

Out[47]: 0.07272346333201705

```
In [36]: # out-of-sample
mape(y_test, model.predict(X_test))
```

```
585/585 [=====] - 2s 4ms/step
```

Out[36]: 0.07269192021798271

Jelikož se data zdají být poměrně lineární, první neuronku jsem zkusil spíše jednoduší, s pouze jednou skrytou vrstvou. Out-of-sample MAPE ~7 procent.

Neuronka s dvěma skrytými vrstvami

- Input layer = 326 neuronů
- Hidden layer = 16 neuronů
- Batch normalizace vstupu
- Dropout 20%
- Hidden layer = 24 neuronů
- Dropout 10%
- Output layer = 1 neuron

Všechny aktivace RELU

Batch size 12

```
In [49]: # Vytvoření a kompilace modelu
model2 = Sequential()
model2.add(Dense(16, input_shape=(326,), activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.2))
model2.add(Dense(24, activation='relu'))
model2.add(Dropout(0.1))
model2.add(Dense(1, activation='relu'))

model2.compile(loss='mean_squared_error', optimizer='adam',
               metrics='mean_absolute_percentage_error')

model2.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 16)	5232
batch_normalization_3 (Batch Normalization)	(None, 16)	64
dropout_5 (Dropout)	(None, 16)	0
dense_10 (Dense)	(None, 24)	408
dropout_6 (Dropout)	(None, 24)	0
dense_11 (Dense)	(None, 1)	25
Total params: 5,729		
Trainable params: 5,697		
Non-trainable params: 32		

```
In [53]: model2.fit(X_train, y_train,  
                  epochs=20, batch_size=12,  
                  verbose=1)
```

Epoch 1/20
6233/6233 [=====] - 15s 2ms/step - loss: 11434965
6.0000 - mean_absolute_percentage_error: 16.2468

Epoch 2/20
6233/6233 [=====] - 15s 2ms/step - loss: 11358037
6.0000 - mean_absolute_percentage_error: 16.2197

Epoch 3/20
6233/6233 [=====] - 14s 2ms/step - loss: 11299057
6.0000 - mean_absolute_percentage_error: 16.1990

Epoch 4/20
6233/6233 [=====] - 14s 2ms/step - loss: 11185695
2.0000 - mean_absolute_percentage_error: 16.1351

Epoch 5/20
6233/6233 [=====] - 14s 2ms/step - loss: 11302000
8.0000 - mean_absolute_percentage_error: 16.1746

Epoch 6/20
6233/6233 [=====] - 14s 2ms/step - loss: 11216140
8.0000 - mean_absolute_percentage_error: 16.1497

Epoch 7/20
6233/6233 [=====] - 15s 2ms/step - loss: 11182761
6.0000 - mean_absolute_percentage_error: 16.1415

Epoch 8/20
6233/6233 [=====] - 15s 2ms/step - loss: 11182588
8.0000 - mean_absolute_percentage_error: 16.0981

Epoch 9/20
6233/6233 [=====] - 15s 2ms/step - loss: 11038562
4.0000 - mean_absolute_percentage_error: 16.0056

Epoch 10/20
6233/6233 [=====] - 15s 2ms/step - loss: 11109351
2.0000 - mean_absolute_percentage_error: 16.0295

Epoch 11/20
6233/6233 [=====] - 14s 2ms/step - loss: 11070683
2.0000 - mean_absolute_percentage_error: 16.0019

Epoch 12/20
6233/6233 [=====] - 14s 2ms/step - loss: 11070185
6.0000 - mean_absolute_percentage_error: 15.9697

Epoch 13/20
6233/6233 [=====] - 15s 2ms/step - loss: 11051803
2.0000 - mean_absolute_percentage_error: 15.9662

Epoch 14/20
6233/6233 [=====] - 15s 2ms/step - loss: 11008112
8.0000 - mean_absolute_percentage_error: 15.9082

Epoch 15/20
6233/6233 [=====] - 14s 2ms/step - loss: 10973395
2.0000 - mean_absolute_percentage_error: 15.8726

Epoch 16/20
6233/6233 [=====] - 14s 2ms/step - loss: 10925093
6.0000 - mean_absolute_percentage_error: 15.9058

Epoch 17/20
6233/6233 [=====] - 14s 2ms/step - loss: 10939073
6.0000 - mean_absolute_percentage_error: 15.8735

Epoch 18/20
6233/6233 [=====] - 15s 2ms/step - loss: 10936861
6.0000 - mean_absolute_percentage_error: 15.8905

Epoch 19/20
6233/6233 [=====] - 14s 2ms/step - loss: 10802593

```
6.0000 - mean_absolute_percentage_error: 15.8455  
Epoch 20/20  
6233/6233 [=====] - 15s 2ms/step - loss: 10831861  
6.0000 - mean_absolute_percentage_error: 15.8034
```

```
Out [53]: <keras.callbacks.History at 0x7f75390238e0>
```

```
In [54]: mape(y_test, model2.predict(X_test))
```

```
585/585 [=====] - 1s 2ms/step
```

```
Out [54]: 0.0732175049171604
```

Zkusil jsem trochu slozitejsi neuronku s pridanou druhou skrytou vrstvou. Prislo mi, ze se ztratova funkce snizovala pomalu, zkusil jsem tedy snizit velikost davky, coz prilis nepomohlo. Vysledky nejsou o moc lepsi ani horsi nez jednoduchussi neuronova sit.

Zaver

Myslim, ze se ukazalo, ze ceny letenek jsou dobre predikovatelne a to i pomoci jednoduchussich modelu. Oproti puvodni benchmark linearni regresi performovali vsechny dalsi modely vyborne. Krome `class` jsem pouzil vsechny sloupce, vcetne kategoricalnich pomoci one hot encodingu.

Jako metriku chyby jsem zvolil MAPE (Mean Absolute Percentage Error), tedy prumernou absolutni procentualni chybu. Vysledna chyba se da interpretovat jako prumerna procentualni odchylka predikce od skutecnych hodnot.

Vyber nejlepsiho modelu

Jelikoz dobrych predikcnich schopnosti dosahovaly i jednoduchussi modely, myslim, ze je vhodne vybrat je. A to primarne z duvodu jejich snadne uchopitelnosti i rychlosti treninku.

Naprosto nejlepsich vysledku dosahly regresni stromy za pouziti bagging metody (bootstrap + aggregate). MAPE tohoto modelu byla pouze ~2.8%, coz je fantasticka hodnota.

I standardni regresni strom, nebo model random forest dosahl dobrych vysledku a MAPE ~4.2% a ~4%. U jednoho stromu mam podezreni na over-fitting, u random forestu si myslim, ze jsem nenalezl optimalni parametry.

Dalsim vhodnym modelem by mohla byt linearni regrese regularizovana pomoci elastic net. Pri pouziti L1 (Lasso) regularizace dosahl tento model out-of-sample chyby asi 11%, coz na linearni regresi neni vubec spatne. Hledani optimalnich parametru vsak trvalo pomerne dlouho.

Myslim, ze v pripade aktualnich dat nema vyznam vuzivat neuronove site. Ackoliv i jednoduchucha sit mela vcelku dobre predikcni schopnosti, neni lepsi nez rozhodovaci

stromy a její trénink a interpretace je složitější. Je možné, že při nalezení lepších hyperparametru, by mohla neuronka dosáhnout stejné kvality predikce jako stromy. Aktivací funkce všech neuronů jsem vybral RELU.

K Support Vector modelům se radši moc nevyjadřuji :D Myslím, že by mělo být možné je zlepšit. Jejich trénování trvalo ale výrazně nejdelší ze všech modelů v tomto notebooku, takže jsem se jim moc nevěnoval.

Jakou míru chyby je možné očekávat?

Rekl bych, že očekatelná míra chyby je rovna hodnotám zmíněným výše. Všechny uvedené MAPE odchylky jsou počítány out-of-sample, tedy na validacích 20% původních dat, které nebyly použity pro trénink. Díky většímu rozsahu dat bylo dost pozorováno pro trénink i pro testování, a nezaznamenal jsem výrazný overfitting.